



Università degli Studi di Milano Bicocca

Department of Informatics, Systems and Communication DISCo

Master's degree in Information Technology

## **Definition and evaluation of a strategy for Real-Time Traffic Engineering in SD-WAN**

**Supervisor: Marco Savi, PhD**

**Thesis by:**  
Michele Beccari  
856608

**Academic Year 2023-2024**

## Abstract

Wide Area Networks (*WANs*) are adopted by large companies and organizations to overcome the potential reliability issues that may arise by just using general connectivity (i.e the standard connections offered by Internet Service Providers (*ISPs*), usually sold for personal use).

However, while offering greater guarantees, setting up a full WAN connectivity is massively more expensive than using ISPs; some WANs may even require renting actual cables from ISPs. SD-WANs (*Software Defined WAN*) offer an hybrid approach to find a middle ground between the reliability of standard WANs and the economic advantage of general connectivity.

When using both standard WAN connectivity solutions and general connectivity traffic may be routed through the different interfaces according to the requirements of the source of the traffic itself. (i.e the more expensive interfaces are only used when strictly necessary)

This paper tries to explore a strategy to route traffic across different kinds of interfaces to minimize the overall transmission cost, while still respecting the requirements of each source.

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Types of connectivity . . . . .	1
1.2	Software defined networking and SD-WAN . . . . .	2
1.2.1	Networking devices . . . . .	2
1.2.2	Definition . . . . .	2
1.2.3	Openflow . . . . .	2
1.2.4	SD-WAN . . . . .	3
<b>2</b>	<b>Problem definition</b>	<b>5</b>
2.1	Load balancing as a linear optimization problem . . . . .	5
2.1.1	Extending the problem . . . . .	6
2.2	An existing strategy . . . . .	7
2.2.1	Intent-Based Routing Policy . . . . .	7
2.2.2	The linear optimization problem . . . . .	7
2.2.3	Estimating the delay . . . . .	9
2.2.4	Estimating the error rate . . . . .	9
<b>3</b>	<b>Implementation</b>	<b>10</b>
3.1	Emulating a network of computers . . . . .	10
3.1.1	Implementing the simulation . . . . .	10

# Introduction

## 1.1 Types of connectivity

The types of internet connections can be usually divided in two main categories, **Generalized connectivity** and **Dedicated connectivity**.

### Generalized connectivity

The most popular kind of internet connection provided by an *ISP* (*Internet Service Provider*) , usually adopted by residential users.

It is based on an hierarchical structure where a set of autonomous system (*AS*) communicate between each other: any two user adopting this type of connection are able to communicate using different protocols (*IGP* protocol if the users are in the same *AS* and *EGP* otherwise)

### Dedicated connectivity (WAN)

A specialized kind of connection, mainly used by businesses (i.e to connect the different points of sale of a large company or the different branches of a bank).

Organizations usually opt in into this kind of connection to ensure higher standards in regards to different metrics (i.e lower delays, higher availability) when compared to the *best-effort* approach of generalized connectivity.

Dedicated connectivity is still offered by internet service providers, but it is several orders of magnitude more expensive than generalized connectivity.

A network adopting this approach is called a *WAN* (Wide Area Network); A WAN is a network of LANs (Local Area Network, i.e the branches) where the distance between the LANs is usually in the tens or hundreds of kilometers range.

By using a *border router* in each LAN, it is possible to communicate with the other LANs in the WAN. The LANs can be connected in different ways:

1. *Dedicated physical WAN* : the business buys the full physical infrastructure that the networks is built on (the cables, the switches etc.).  
While the company is the owner of the whole WAN, the cost is extremely high and network management (e.g enforcing security) is the the company's responsibility.
2. *Leased Lines* : the business relies on an ISP to establish private circuits between the branches. (e.g branch 1 must be connected to branch 2, which must be connected to branches 3 and 4)  
The private circuit may be created by assigning a dedicated wavelength to a particular customer. This solution offers only partial control, but is cheaper and offsets some of the network's management issues to the ISP.
3. *MPLS networks*: The business still relies on an ISP to provide connectivity, but unlike leased lines there are no established circuits between the branches. Instead, a mesh connectivity is provided , with a certain QoS guarantee.  
For each branch a *customer edge router* is installed, allowing the branch to access the MPLS network via a *provider edge router*.  
The *provider edge router* (PE) is in between the MPLS network and the companies, one PE may manage connections for multiple customers.
4. *SD-WAN*: this kind of WANs will be discussed in detail in the next section.

## 1.2 Software defined networking and SD-WAN

### 1.2.1 Networking devices

A typical network device (e.g routers, firewalls...) is usually composed by two logical components:

1. A data plane
2. A control plane

The data plane is used to handle individual data packets locally. For example the data plane in a router is used to determine to which interface a packet should be forwarded to. The data plane usually applies a large amount of small, fast operations.

The control plane is used to handle control messages from other devices in the network, and allows the network device to work from a network wide perspective.

The messages from other devices may be used to configure policies which are then applied by the data plane (e.g the control plane receives a policy that determines which packets should be filtered, while the data plane does the actual filtering).

Compared to the data plane, the control plane is used for more complex operations that usually require coordination along network devices.

### 1.2.2 Definition

The idea behind software defined networking (SDN) is to de-couple the control plane from the data plane, not only on a logical level but also physically.

The control plane is no longer part of the network devices and is instead moved to a centralized software component called a *controller*.

The basic architecture of a SDN is composed by the controller and by a series of network devices which in this context are called *switches*.

A SDN controller

1. Is able to communicate with the switches
2. Can be programmed: it is now possible to write software that defines how the whole network behaves.  
The switches keep executing their data plane locally, according to the policies received by the controller.
3. Is able to acquire information about the network it is managing, e.g it can discover the topology of the network to know where the switches are placed and how they are connected.

The goal of the controller is to communicate with all the switches to inject into them the desired network behavior

The controller is logically placed between the network infrastructure and the application layer.

Using a series of high level programmable interfaces (also called *APIs Application Programmable Interfaces*) the controller is able to communicate with its neighboring layers.

Using the APIs connected with to the application layer (the *northbound*) interfaces various network devices can be implemented, e.g a network balancer, a firewall.

The communication between the controller and the switches is made possible by a series of *southbound* interfaces.

Using these interfaces the controller can effectively apply the policies to the network, e.g it can send traffic rules to the switches connected to the controller.

Inside each switch there is a set of *generic tables* (also called *flow tables*) that are used to store the configuration sent by the southbound interfaces.

The southbound interfaces can also be used in "reverse": the switches are able to communicate with the controller using a series of predefined messages.

### 1.2.3 Openflow

Despite being relatively new, software defined networking became a popular approach for organizing new networks (e.g in 5G networks) Compared to other widespread network technologies such as the TCP/IP stack, software defined networking is a relatively new approach: the first paper defining protocol for SDN (the *Openflow Protocol*) was published in 2008.

The OpenFlow protocol defines a sets of API for remote communication with network plane devices, i.e. a southbound interface.

SDN switches inside the context of OpenFlow are also called *generic* switches; This is to avoid confusion with the usual meaning of the term switch, which usually refers to a device used to manage level 2 functionalities (i.e. the data link level, used to handle the physical addressing with MAC addresses). Openflow switches are able to interact up to the fourth layer (i.e. the transport layer, which may include protocols such as TCP or UDP).

A large portion of an OpenFlow Switch is composed by functionalities related to the data plane, the only portion dedicated to the control plane is delegated to the *control channel*.

The control channel includes one or more OpenFlow channel that allows the switches to communicate with one or more controllers (multiple controllers may be used to avoid a single point of failure in the infrastructure).

OpenFlow switches include three kind of tables, used for handling data plane functionalities:

- *Flow tables*: the most important kind of table, they can be chained in a pipeline. They contain the rules that are applied to the incoming and outgoing traffic.
- *Group tables*: Users of a SDN can be assigned to different group; group tables allow traffic to be sent in multicast to users of a specific group (e.g. group tables can be used to send outgoing traffic to multiple ports if necessary)
- *Meter table*: table used to collect statistic about the current network. The gathered data can then be used to handle QoS (e.g rate limiting for traffic directed to a specific port)

### Configuring the tables

To work properly, the flow tables inside the switches have to be configured by the controllers in the network. The solution is usually an hybrid of two types of configuration:

1. *Proactive configuration*: When the switch is first inserted into the network, the controller proactively fills the tables.  
The entries in the tables are static and do not change across the lifespan of the switch in the network.  
This kind of approach is obviously unfit for rapid changing scenarios such as those encountered in managing networks.
2. *Reactive configuration*: Initially there no entries in the flow tables; the tables are fully configured at runtime according to the network state.  
If a switch does not know where a package should be sent to, it forwards the packet to the controller.  
The controller is aware of the full topology of the network, and is then able to edit the received package to include the routing information, before sending the package back to the switch..  
This leads to increased round trip time for the first packets sent across the network.  
However, communication between the switches and controller usually represent a bottleneck for the network's performance; a large amount of packets sent to and from the controller may cause congestion in the network and limit the available bandwidth.

Usually the preferred solution is based on an hybrid approach.

A starting set of rules is used at the start up of the network, to cover the expected data flows in the network.

For any unknown data flow that may arise during the lifespan of the network new routing rules will be issued by the controller.

## 1.2.4 SD-WAN

### SD-WAN

The principles of software defined networking can be applied when building a WAN; While SD-WAN were introduced recently (in 2014) their popularity increased due to the possible decrease in costs when compared to traditional WAN solutions.

SD-WAN works by stipulating multiple contracts of generalized internet connections (e.g 4G, POM ...) and then combining the different networks to guarantee a certain quality of service level.

Like the MPLS approach, a CPE (*Customer Premises Equipment*) is installed in each branch of the business.

The CPEs are connected to the branch's LAN and to multiple generalized connectivity networks.

The traffic generated by the LAN may then travel in different networks, potentially even from different ISPs.

An SDN controller, in this specific context called an *SD-WAN Controller* is responsible to communicate to each CPE how the outgoing traffic should be managed (e.g. two branches may be connected using a fiber cable, if the cable connection fails the controller will instruct the branches to communicate using a 4G network).

MPLS may still be used in a SD-WAN context, a company could decide to use other (usually cheaper) kind of connections for most use cases and reserve a MPLS connection only exceptionally.

### **Load balancing and SLA in SD-WANs**

As we previously discussed, in SD-WANs a controller's task is to instruct CPEs on how they should distribute outgoing traffic among the available links.

The traffic handled in a SD-WAN context is often heterogeneous and generated by different applications (i.e. VoIP, emails, video streaming ...).

Different applications have usually varying requirements in regards to the acceptable delay and the packet loss that can happen during the transmission of data.

Two important terms have to be introduced to describe the needs of each network application:

1. *QoS (Quality of Service)*: a quality indicator used to measure the service level of a given network communication. It is defined with a series of guarantees that must be respected during a network communication.  
In this paper QoS will always be defined in absolute terms (e.g. end to end latency is less than 5 milliseconds, 2% of packets have been dropped), however it can also be defined in relative terms (i.e. how a given traffic flow is treated with respect to the others)
2. *SLA (Service Level Agreement)*: the SLA specifies the QoS that has to be guaranteed for a given application. It provides a series of bound on various measurable parameters (i.e. the latency *must* be less than 5 milliseconds)

## Problem definition

### 2.1 Load balancing as a linear optimization problem

One of the goals of SD-WAN is to minimize costs while still maintaining a set of requirements for a given set of applications.

With *application* we mean any kind of source of internet traffic with a defined set of requirements. An application has a *life cycle*: it starts, it generates data according to arbitrary rules and it terminates when no more data has to be sent.

As a first step in formalizing the problem, we can picture a scenario in which we have a CPE sending an arbitrary amount of packets to a sink.

The CPE is connected to the sink through a set of distinct links.

The goal of the CPE is to allow an *application* to transmit an arbitrary amount of packets to the sink. The transmission of data has some bounds in regards to:

- Average bandwidth
- Average delay
- Average Packet loss

An important assumption on the application requirements is that they only apply to the average: if an application requires the average delay to be 100ms some of the packets can be sent with an higher delay, as long as the overall delay stays under the desired threshold.

Managing applications that have stricter requirements (e.g all packets must always be sent with a maximum a delay of 100ms) go beyond the scope of our strategy.

Each link has different characteristics in regard to:

- Available bandwidth (*i.e how many bytes can be sent per second through the link*)
- Delay (*i.e how long it takes for a packet to reach the sink*)
- Cost (*i.e a metric used to track how convenient is to use a link instead of another, the section 2.1 goes into further details*)
- Packet loss (*i.e how many packets are lost and/or need to be sent multiple times?*)

#### Measuring the cost of links

Each link has a cost associated with sending a packet on the link itself.

The cost of sending a packet is not a monetary cost, rather it is a measure of how convenient it is to use a given link.

Monetary costs for different kinds of connections are measured in multiple ways: i.e MPLS connections charge increasing amount of money depending on the amount of traffic, while typical consumer ISPs charge a fixed amount regardless of the amount of traffic.

This cost for each packet is fixed for each link, and it does not change through the lifetime of any application.

This means that when assigning the packet cost to a given link we assign lower costs to the links that have a fixed monetary cost, and an higher cost to the link the have a pay-as-you-go model.



## The problem as a linear optimization problem

We can formalize the problem as a linear optimization problem; As a starting point we consider an application sending an amount  $N$  of packets of fixed size using three different links.

The amount of packets is fixed and generated by a single application: our goal is to minimize the overall cost of the transmission while respecting the different requirements.

To simplify the problem we set the bandwidth of the links in a way that it is always a multiple of the size of a single packet (i.e the bandwidth of a link can only be an integer, 1 packet per second, 5 packet per second ...)

We consider a set of links  $L$ , to visualize the example we can set  $L = \{l_1, l_2, l_3\}$ .

For each one we define

- The link delay:  $d_1, d_2, d_3$ , measured in *ms*
- The link bandwidth:  $b_1, b_2, b_3$ , measured in *packets per second (pps)*
- The cost for sending a single packet  $c_1, c_2, c_3$
- The probability of losing a packet, expressed as a number in the range  $[0, 1)$

The total cost time required to send the packets can be expressed as:

$$N_1 \cdot c_1 + N_2 \cdot c_2 + N_3 \cdot c_3$$

Where  $N_i$  is the amount of packets sent on the  $i$ th interface.

We also have to consider the bound on average delay:

$$\frac{N_1 \cdot d_1 + N_2 \cdot d_2 + N_3 \cdot d_3}{N} \leq \text{delay}$$

The bound on average bandwidth:

$$\frac{N_1 \cdot b_1 + N_2 \cdot b_2 + N_3 \cdot b_3}{N} \leq \text{bandwidth}$$

The bound on packet loss, expressed as a percentage:

$$\frac{N_1 \cdot p_1 + N_2 \cdot p_2 + N_3 \cdot p_3}{N} \leq \text{packet loss}$$

Finally, we need to express a bound to make sure every packet is actually sent.

$$N_1 + N_2 + N_3 \geq N$$

This simplified version of the problem can be expressed as:

$$\left\{ \begin{array}{ll} \min & N_1 \cdot c_1 + N_2 \cdot c_2 + N_3 \cdot c_3 \\ \text{subject to:} & \frac{N_1 \cdot d_1 + N_2 \cdot d_2 + N_3 \cdot d_3}{N} \leq \text{delay} \\ & \frac{N_1 \cdot b_1 + N_2 \cdot b_2 + N_3 \cdot b_3}{N} \leq \text{bandwidth} \\ & \frac{N_1 \cdot p_1 + N_2 \cdot p_2 + N_3 \cdot p_3}{N} \leq \text{packet loss} \\ & N_1 + N_2 + N_3 \geq N \end{array} \right. \quad (2.1)$$

### 2.1.1 Extending the problem

While the linear optimization problem allows for an easy and clear solution, it can only work if some assumptions are made:

- Only a single application is sending data at any given time
- The amount of traffic generated by the application must be known beforehand (this is usually unrealistic, unless for specific situations, i.e file transfer)

This linear approach does also not take into account delays introduced by the outgoing traffic; the theoretical delay on a given link is only a lower bound.

When sending multiple packages over an interface the time required to send the packages that are already into the queue must also be taken into account (e.g if a package is sent through an interface where  $n$  packages are already queued, the  $n$  packages must be sent first)

The linear optimization strategy is still useful to set a lower bound on the lowest possible overall transmission cost.

To allow for modeling of more complex scenarios, we extend the problem definition to allow the presence of multiple applications.

In this new scenario, each application can now generate traffic at any point of its lifetime: the total amount of traffic can not be determined until the end of the application lifecycle.

Examples of different applications may include some that generates:

- A fixed amount of traffic every second.
- A random amount of traffic every hundred milliseconds
- A burst of traffic when starting the application, then no more.
- Any kind of commonly studied traffic model (i.e Poisson distribution traffic model, Pareto distribution model)
- A simplified model of traffic (i.e sinusoidal traffic)

Any proposed solution must:

- Be able to account for multiple sources of traffic
- Be able to account for the delay caused by the queuing of packages on a given interface

## 2.2 An existing strategy

### 2.2.1 Intent-Based Routing Policy

In the paper *Intent-Based Routing Policy Optimization in SD-WAN* [1] a strategy is discussed to optimize different metrics in a scenario with multiple applications (called "flow groups" in the original paper) with different QoS requirements; the traffic generated from a set of applications has to be distributed on a set of available links.

The goal of the strategy is to provide the "best" possible split ratio for the traffic generated by the different applications among the available links.

What is the "best" split ratio can be arbitrarily decided: in the original paper multiple metrics are proposed such as minimizing link utilization, or even improving the quality of the transmission beyond what is strictly required.

The idea behind the strategy is run a linear optimization problem periodically so that it can adjust to changing demands by the different applications.

The linear optimization resolution relies on an estimate to determine what the delay on a given link for a given application will be in the future; a similar approach could be used to estimate an error rate for a given application on a given link, as a starting point thought we suppose the error rate is always fixed given an link.

### 2.2.2 The linear optimization problem

The original paper provides a formulation with a linear programming model that may include different objectives, including:

1. An objective that measures if the SLA for a given application (in the original paper called a *flow group*) can be violated. We will not allow for any violation of SLA, so this objective can be safely ignored.
2. An objective function that minimizes the MLU (Maximum link utilization)
3. An objective function that can eventually be used to decrease the delay beyond SLA requirements.
4. An objective function that may be used to minimize costs.

For our purposes, we can modify the proposed model. While the importance of the various objective functions in the original formulation can be adjusted using weights, we can ignore most of the proposed functions (i.e setting their weight to 0).

We consider a SD-WAN composed by a CPE sending data to a sink; the network is composed by:

- A set of links  $L$  with bandwidth  $B_l$ ,  $\forall l \in L$
- A set of applications,  $A$ . For each application  $a \in A$ , the traffic demand is denoted with  $t_a$ , the required delay over all links is  $\bar{D}_a$  and the required error rate over all links is  $\bar{E}_a$

The modified formulation is:

$$\min \sum_{l \in L} c_l \cdot LU_l \quad (1)$$

$$\text{s.t. } LU_l = \sum_{a \in A} t_a x_l^a \leq C_l, \quad \forall l \in L, \quad (2)$$

$$f_l^a(x) \leq D_a, \quad \forall a \in A, \forall l \in L, \quad (3)$$

$$D_a \leq \bar{D}_a, \quad \forall a \in A, \quad (4)$$

$$e_l^a(x) \leq E_a, \quad \forall a \in A, \forall l \in L, \quad (5)$$

$$E_a \leq \bar{E}_a, \quad \forall a \in A, \quad (6)$$

$$\sum_{l \in L} x_l^a = 1, \quad \forall a \in A \quad (7)$$

The variables are:

Variable	Range	Notes
$c$	$R$	The cost coefficient for a link
$x_l^a$	$[0, 1]$	The split ratio of each application $a \in A$ over each link $l \in L$
$LU_l$	$[0, 1]$	The utilization in percentage for each link $l \in L$
$D_a$	$R$	The maximum delay estimated by application $a$ over all overlay link
$f_l^a(x)$	$R$	The delay function that provides the delay of application $a$ over overlay link $l$ considering the assignment given by $x$ . This function will be discussed in greater detail in the 2.2.3 section.
$e_l^a(x)$	$R$	The error rate function that provides the error rate of application $a$ over overlay link $l$ considering the assignment given by $x$ . This function will be discussed in greater detail in the 2.2.4 section

The constraints are:

Constraint	Meaning
(2)	Ensures that the capacity of each link is satisfied
(3)	Computes the delay for each application over each link
(4)	Verifies that satisfaction of SLA delay requirements
(5)	Computes the error rate for each application over each link
(6)	Verifies that satisfaction of SLA error rate requirements
(7)	Ensures that all pending packets are sent

### 2.2.3 Estimating the delay

The most intuitive approach for solving a linear optimization problem would be to keep using the cheapest links that meet the requirements to send the incoming packets.

The problem with this approach is that overusing an interface will inevitably lead to congestion.

While we know the upper bound of the bandwidth of any given link (i.e  $B_l, \forall l \in L$ ), we must also take into account the reduction of bandwidth caused by the packets we previously sent.

In this paper we will rely on a very basic delay function, that it is purely based on how many packets arrived on a link up to a point in time.

When estimating the delay of sending  $n$  packets on a given link, we can rely on the arrival rate up to that instant of time to approximate how much load is on a current link.

$$D_a = \frac{n}{\text{arrival rate}}$$

Where the arrival rate can be computed by considering how many packets were sent through the link up to a given point in time:

$$\text{arrival rate} = \frac{\text{total amount of packets}}{\text{current time}}$$

Estimating the delay constitutes a challenge that opens for a lot of possible extensions to the current implementation, this will be further discussed in the future work session.

### 2.2.4 Estimating the error rate

Each link has a set error rate, expressed as a percentage, that affects how many packets actually reach the destination. For the scope of this paper the error rate will be fixed for every application and for every link.

However the error rate can potentially be any user-defined function to account for possible changes in the links (e.g the error rate on a given link could rise when incoming traffic goes over a certain threshold)

# Implementation

## 3.1 Emulating a network of computers

To test the strategy we defined we setup a simple network with two nodes:

1. A source node representing the LAN adopting a SD-WAN connectivity
2. An arbitrary node to which the data is sent: in a real life scenario it would be another LAN (e.g another branch); in our case it will be a sink node used to collect data about the strategy (i.e how many packets were received, from which application they were generated)

The two nodes will then have an arbitrary amount of links connecting them: each link represent an element  $l \in L$  of.

Once the nodes are connected we have to define the application set  $A$ . Each application is a piece of code that may generate any amount of packets at any point in time, until it finally stops.

At a very high level the code for the simulation is:

---

**Algorithm 1** The simulation algorithm

---

```
while ShouldBeRunning do  
    Result  $\leftarrow$  executeStrategy( $A, L$ )  
    enqueuePackages(Result)  
    if there are no longer pending packages and all applications have stopped then  
        ShouldBeRunning  $\leftarrow$  false  
    end if  
end while
```

---

When running the experiments, the applications will stop after having generated a pre-determined amount of traffic. The applications will run (and therefore generate traffic) for a set amount of time and the simulation will stop when all the packets will have been sent.

### 3.1.1 Implementing the simulation

The simulations will be implemented using the C++ programming language, the ns3 simulator and the or-tools library.

#### Discrete event simulation

Using a discrete-event simulation we can model all events in our system.

With event we generally mean any piece of code that has to be ran at a specific time and alters the state of the emulated system (i.e alters the state of the network).

The simulator keeps track of when each event should be executed and automatically updates the state of system accordingly.

For example, if the event is *send packet on link  $l_1$  to the sink*, the simulator updates the state of the system with the information that there is a packet being sent on  $l_1$ .

The system will then react to the new state: to continue on with our example after the necessary time for the packet to traverse the link will have passed the state of the system will be updated with the information that the sink has received a packet.

It is important to note that the time used by the simulator is more often than not unrelated to actual physical time: we may want to emulate a really slow connection that takes up seconds to send a packet but the actual time to run the simulation will most likely be in milliseconds.

The opposite could also be true, we may want to emulate a really fast connection were gigabytes of data are sent through a link.

We could potentially then implement a sink that prints all the received data and takes minutes to do so, but no *simulator* time will have passed (printing data is something that is outside the bounds of the simulation)

Discrete event simulator are usually built to keep an internal timer based only on the events that alter the state of the simulated system.

### NS3

One very popular choice for network discrete-event simulation in networks is *NS3* [2].

NS3 is a discrete-event simulator written with the C++ programming language which allows for easy simulation of networks. Most components that constitute a network can easily be emulated with NS3: in our case we can setup two nodes and connect them using an arbitrary amount of links.

All the technical details such as setting up the routing, defining the error rate of a given link, the IP addresses etc... can be handled with the tools offered by NS3.

We will be using NS3 for all the experiments: while the setup for our simulation is tailored to meet our use case, it can easily be modified for future works (e.g we use a static percentage-based error rate for our links, but NS3 allows defining any arbitrary error model that can then be applied to a link)

### OR Tools

As we previously stated, the core of the strategy is solving a linear optimization problem.

A popular tools for handling linear optimization problem is *OR-Tools* [3]

. According to its creators, OR-Tools is a *an open-source, fast and portable software suite for solving combinatorial optimization problems..* Using OR-Tools we can offload the task of solving the linear optimization problem.

### CMake

Such as with most moderately complex projects, code from multiple sources has to be organized and then compiled together.

Unlike other programming languages, C++ does not have an official tool for defining projects and their dependencies.

One of the most popular software build system is *CMake* [2]: we will rely on CMake to combine the custom code (i.e the strategy) and the dependencies (NS3 and or-tools).

## Bibliography

- [1] Pham Tran Anh Quang et al. Intent-based routing policy optimization in sd-wan. 2022.
- [2] Ns3. <https://www.nsnam.org/>, last accessed on 04/03/2025.
- [3] Or-tools. <https://developers.google.com/optimization?hl=it>, last accessed on 04/03/2025.